

# 1 BM3DORNL: High-Performance BM3D Denoising for 2 Neutron Tomography

3 Chen Zhang  <sup>1¶</sup>, Jean-Christophe Bilheux  <sup>2</sup>, Dmitry Ganyushin<sup>1</sup>, and Pete  
4 Peterson  <sup>1</sup>

5 1 Computing and Computational Sciences Directorate, Oak Ridge National Laboratory, Oak Ridge, TN,  
6 USA 2 Neutron Sciences Directorate, Oak Ridge National Laboratory, Oak Ridge, TN, USA ¶  
7 Corresponding author

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

## Software

- [Review](#) ↗
- [Repository](#) ↗
- [Archive](#) ↗

Editor: [Open Journals](#) ↗

## Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

## License

Authors of papers retain copyright  
and release the work under a  
Creative Commons Attribution 4.0  
International License ([CC BY 4.0](#))<sup>3</sup>

8 **Notice:** This manuscript has been authored by UT-Battelle, LLC, under contract DE-AC05-  
9 00OR22725 with the US Department of Energy (DOE). The US government retains and the  
10 publisher, by accepting the article for publication, acknowledges that the US government  
11 retains a nonexclusive, paid-up, irrevocable, worldwide license to publish or reproduce the  
12 published form of this manuscript, or allow others to do so, for US government purposes. DOE  
13 will provide public access to these results of federally sponsored research in accordance with  
14 the DOE Public Access Plan (<https://www.energy.gov/doe-public-access-plan>).

## 15 Summary

16 BM3DORNL is a high-performance Python library for denoising neutron and X-ray tomography  
17 data using a modified Block-Matching and 3D Filtering (BM3D) algorithm. The library  
18 provides two denoising modes: a generic mode for standard noise removal, and a specialized  
19 streak mode optimized for removing vertical streak patterns in sinograms that manifest as ring  
20 artifacts in reconstructed images. Built with a Rust backend and Python bindings via PyO3,  
21 BM3DORNL achieves over 500× faster processing than comparable BM3D implementations  
22 while maintaining cross-platform compatibility including native Apple Silicon support.

## 23 Statement of Need

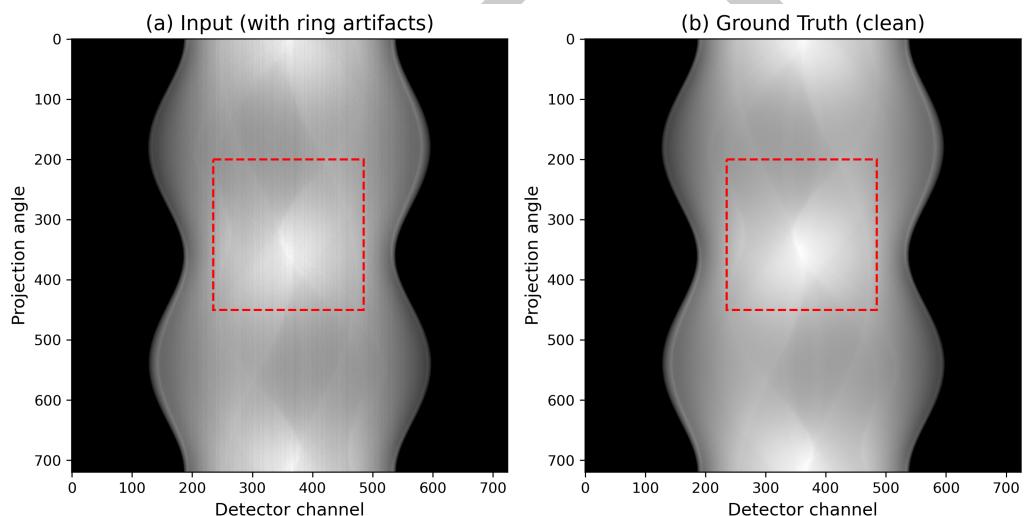
24 Ring artifacts are a persistent challenge in neutron and X-ray computed tomography, arising  
25 from variations in detector pixel response, beam intensity fluctuations, and systematic errors  
26 ([Münch et al., 2009](#)). These artifacts appear as concentric rings in reconstructed images,  
27 obscuring sample structure and degrading quantitative analysis. Pre-reconstruction streak  
28 removal addresses these artifacts directly in sinograms, avoiding the spatial spreading that  
29 occurs during reconstruction. Existing approaches include wavelet-Fourier filtering ([Münch et  
30 al., 2009](#)), polynomial fitting ([Vo et al., 2018](#)), and sorting-based methods ([Miqueles et al.,  
31 2014](#)), with implementations available in TomoPy ([Gürsoy et al., 2014](#)).

32 Mäkinen et al. ([Mäkinen et al., 2021](#)) demonstrated that applying BM3D ([Dabov et al., 2007](#))  
33 across multiple scales achieves superior artifact removal by exploiting the self-similarity of streak  
34 patterns. However, their bm3d-streak-removal implementation is closed-source, restricted to  
35 non-commercial use, and provides only x86\_64 binaries incompatible with Apple Silicon and  
36 modern Python versions. BM3DORNL provides the neutron imaging community with an open-  
37 source, MIT-licensed, high-performance implementation that enables both high-throughput  
38 batch processing and interactive parameter tuning.

## 39 State of the Field

40 Several software packages implement pre-processing streak removal for tomography: TomoPy  
 41 ([Gürsoy et al., 2014](#)) provides wavelet-Fourier filtering ([Münch et al., 2009](#)) and Vo's sorting/fitting  
 42 methods ([Vo et al., 2018](#)); ASTRA Toolbox ([Aarle et al., 2016](#)) offers GPU-accelerated  
 43 preprocessing; and bm3d-streak-removal ([Mäkinen et al., 2021](#)) implements multiscale BM3D  
 44 but remains closed-source and platform-limited.

45 [Figure 1](#) shows the benchmark input data: a synthetic sinogram ( $720 \times 725$  pixels) with  
 46 simulated ring artifacts and the corresponding clean ground truth. All benchmark results  
 47 are from Linux x86\_64 to enable comparison with bm3d-streak-removal, which provides only  
 48 x86\_64 binaries. We compared eight methods: four BM3DORNL variants (streak, generic,  
 49 multiscale, and Fourier-SVD), three TomoPy algorithms (wavelet-Fourier, sorting-fitting, and  
 50 sorting-based), and the original bm3d-streak-removal. The multiscale variant implements  
 51 the pyramid approach from Mäkinen et al. ([Mäkinen et al., 2021](#)), while Fourier-SVD is a  
 52 lightweight FFT-guided method that achieves comparable quality at  $50\times$  faster speeds.

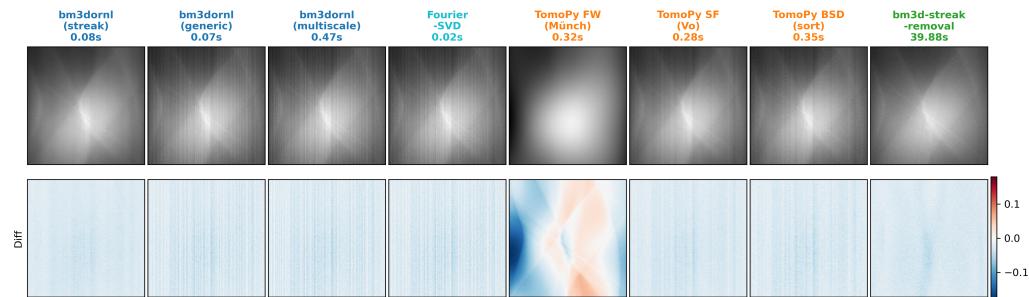


**Figure 1:** Benchmark input data. (a) Input sinogram with simulated ring artifacts. (b) Ground truth (clean sinogram). The dashed rectangle indicates the crop region shown in [Figure 2](#).

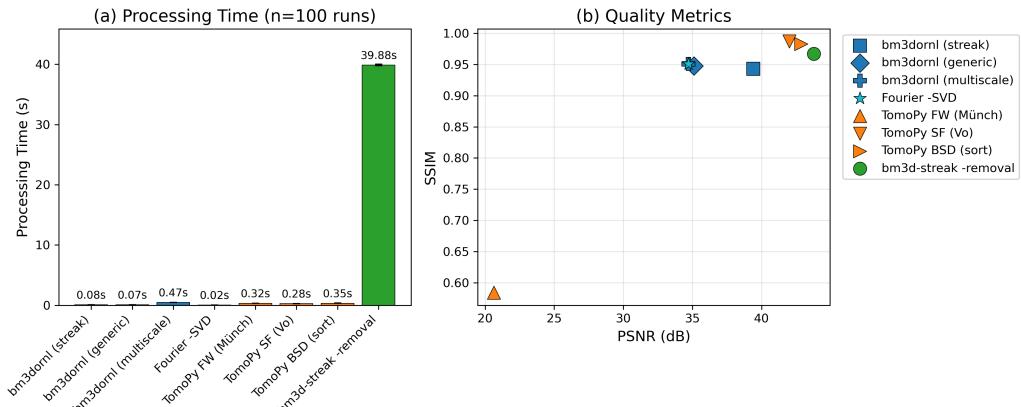
53 **Speed comparison:** [Figure 3\(a\)](#) shows processing times ( $n=100$  runs on Linux x86\_64).  
 54 Fourier-SVD is the fastest method at 0.017 seconds, achieving  $2300\times$  speedup over bm3d-  
 55 streak-removal (39.9 seconds). The BM3D-based methods span a performance range: standard  
 56 BM3DORNL processes sinograms in 0.078 seconds ( $511\times$  faster than bm3d-streak-removal),  
 57 while multiscale BM3DORNL takes 0.475 seconds ( $84\times$  faster) due to pyramid processing.  
 58 TomoPy methods achieve processing times of 0.28–0.35 seconds. Cross-platform performance  
 59 differs substantially: TomoPy runs 2.5–3× slower on Apple Silicon than Linux x86\_64 (likely  
 60 due to unoptimized C extensions), whereas BM3DORNL runs at comparable speeds (within  
 61 10%) across platforms due to its Rust backend.

62 **Quality analysis:** [Figure 2](#) shows cropped results from all eight methods with their difference  
 63 images (result minus ground truth). The benchmark reveals that quality metrics alone do not  
 64 tell the full story. TomoPy SF achieves the highest SSIM (0.987, see [Figure 3\(b\)](#)), but difference  
 65 image analysis shows residual vertical stripes in its output. This occurs because SSIM uses  
 66 local windows ( $7 \times 7$  or  $11 \times 11$  pixels), and narrow vertical stripes affect few pixels per window.  
 67 Notably, all BM3DORNL variants achieve excellent SSIM scores (0.943–0.951), with both  
 68 multiscale BM3DORNL and Fourier-SVD reaching 0.951. Fourier-SVD accomplishes this at  
 69 0.017 seconds, demonstrating that aggressive artifact removal and high-quality reconstruction

70 are achievable without computationally expensive multi-scale processing. The difference images  
 71 reveal that all BM3DORNL variants produce clean vertical patterns, indicating successful  
 72 artifact removal, while TomoPy methods leave faint residual artifacts.



**Figure 2:** Method comparison. Top row: Cropped results from each method with processing times. Bottom row: Difference images (result minus ground truth) using a zero-centered RdBu colormap—blue indicates artifact removal, red indicates added artifacts. Data from Linux x86\_64.



**Figure 3:** Performance metrics (Linux x86\_64, n=100 runs). (a) Processing time comparison on linear scale, showing Fourier-SVD's 2300× speedup and standard BM3DORNL's 511× speedup over bm3d-streak-removal. (b) Quality metrics (PSNR vs SSIM) showing trade-offs between methods. BM3DORNL variants (blue/cyan), TomoPy (orange), bm3d-streak-removal (green).

73 **Platform support:** bm3d-streak-removal is unavailable on Apple Silicon (x86\_64 binaries only),  
 74 incompatible with Python 3.11+, and restricted to non-commercial use. BM3DORNL provides  
 75 native performance on all platforms, supports Python 3.12+, and uses the MIT license for  
 76 unrestricted commercial use.

## 77 Software Design

78 BM3DORNL employs a hybrid Python-Rust architecture: the core algorithm is implemented in  
 79 Rust using the rayon crate for work-stealing parallelism, with Python bindings via PyO3 for  
 80 seamless NumPy integration.

### 81 Key optimizations:

- 82 **Integral image pre-screening:** Before computing expensive patch distances, the block  
 83 matching stage uses integral images to compute mean and norm bounds in O(1) time.  
 84 Patches that fail these bounds are skipped, eliminating approximately 80% of distance  
 85 calculations.

86     ▪ **FFT with SIMD acceleration:** BM3DORNL defaults to FFT-based transforms with  
 87       batched row/column passes and specialized  $8 \times 8$  fast paths. FFT plans are computed  
 88       once and shared across threads via Arc, and a direct-mapped transform cache elimi-  
 89       nates redundant computations. An alternative Walsh-Hadamard path is available for  
 90       multiplication-free processing.

91     ▪ **SIMD-optimized block matching:** The  $8 \times 8$  patch distance computation uses platform-  
 92       specific SIMD instructions with a contiguous-memory fast path, reducing the block  
 93       matching inner loop to near-hardware-limit throughput.

94     ▪ **Tile-owned aggregation:** Sinograms are processed in tiles that fit in L2/L3 cache, with  
 95       SIMD-optimized weight accumulation. Per-worker scratch buffers are reused across  
 96       kernel runs, minimizing memory allocation overhead.

97     The library provides a minimal Python API:

```
from bm3dornl import bm3d_ring_artifact_removal

cleaned = bm3d_ring_artifact_removal(
    sinogram,
    mode="streak",
    sigma=0.0, # default to automatic noise estimation
)
```

98     **GUI application:** BM3DORNL includes a native GUI built with the egui framework for  
 99       Rust, installable via pip install bm3dornl[gui] or Homebrew on macOS. The GUI enables  
 100      interactive parameter tuning, side-by-side comparison with difference visualization, HDF5/TIFF  
 101      file loading with dataset browsing, and real-time processing feedback with a fast-mode toggle.  
 102      At 0.08 seconds per sinogram, scientists can explore parameter space interactively at over 12  
 103      frames per second—something impractical with bm3d-streak-removal's 40-second processing  
 104      time.

## 105     Research Impact

106     BM3DORNL is being integrated into processing pipelines at the VENUS and MARS beamlines  
 107       at Oak Ridge National Laboratory. The library provides multiple performance tiers: Fourier-  
 108       SVD enables real-time parameter exploration at 0.017 seconds per sinogram ( $2300\times$  faster  
 109       than bm3d-streak-removal), standard BM3DORNL offers robust denoising at 0.078 seconds  
 110       ( $511\times$  speedup), and multiscale BM3DORNL provides maximum quality at 0.475 seconds  
 111       ( $84\times$  speedup). For batch processing, a 1000-sinogram dataset that would take 11 hours  
 112       with bm3d-streak-removal completes in 17 seconds (Fourier-SVD), 78 seconds (standard  
 113       BM3DORNL), or 8 minutes (multiscale BM3DORNL).

114     The hybrid Rust-Python architecture demonstrates a modern approach to scientific software  
 115       development: Rust provides memory safety and performance portability (single codebase com-  
 116       piles natively to arm64 and x86\_64), while Python ensures integration with the NumPy/SciPy  
 117       ecosystem. This pattern is increasingly valuable as the scientific computing community  
 118       diversifies hardware platforms.

119     BM3DORNL fills a critical gap: the neutron imaging community needed an open-source,  
 120       actively maintained BM3D implementation that works on modern platforms. The MIT license  
 121       enables unrestricted use at national facilities, and native Apple Silicon support ensures scientists  
 122       can use contemporary hardware for analysis workstations.

## <sup>123</sup> AI Usage Disclosure

<sup>124</sup> Generative AI tools (Claude) were used for code assistance and documentation drafting. All  
<sup>125</sup> AI-generated content was reviewed, tested, and validated by the authors.

## <sup>126</sup> Acknowledgements

<sup>127</sup> This research used resources at the Spallation Neutron Source, a DOE Office of Science User  
<sup>128</sup> Facility operated by Oak Ridge National Laboratory. This work is supported by the U.S.  
<sup>129</sup> Department of Energy under Contract No. DE-AC05-00OR22725.

## <sup>130</sup> References

- <sup>131</sup> Aarle, W. van, Palenstijn, W. J., Cant, J., Janssens, E., Bleichrodt, F., Dabravolski, A.,  
<sup>132</sup> De Beenhouwer, J., Batenburg, K. J., & Sijbers, J. (2016). Fast and flexible x-ray  
<sup>133</sup> tomography using the ASTRA toolbox. *Optics Express*, 24(22), 25129–25147. <https://doi.org/10.1364/OE.24.025129>
- <sup>135</sup> Dabov, K., Foi, A., Katkovnik, V., & Egiazarian, K. (2007). Image denoising by sparse 3-d  
<sup>136</sup> transform-domain collaborative filtering. *IEEE Transactions on Image Processing*, 16(8),  
<sup>137</sup> 2080–2095. <https://doi.org/10.1109/TIP.2007.901238>
- <sup>138</sup> Gürsoy, D., De Carlo, F., Xiao, X., & Jacobsen, C. (2014). TomoPy: A framework for  
<sup>139</sup> the analysis of synchrotron tomographic data. *Journal of Synchrotron Radiation*, 21(5),  
<sup>140</sup> 1188–1193. <https://doi.org/10.1107/S1600577514013939>
- <sup>141</sup> Mäkinen, Y., Marchesini, S., & Foi, A. (2021). Ring artifact reduction via multiscale nonlocal  
<sup>142</sup> collaborative filtering of spatially correlated noise. *Journal of Synchrotron Radiation*, 28(3),  
<sup>143</sup> 876–888. <https://doi.org/10.1107/S1600577521001910>
- <sup>144</sup> Miqueles, E. X., Rinkel, J., O'Dowd, F., & Bermúdez, J. S. V. (2014). Generalized titarenko's  
<sup>145</sup> algorithm for ring artefacts reduction. *Journal of Synchrotron Radiation*, 21(6), 1333–1346.  
<sup>146</sup> <https://doi.org/10.1107/S1600577514016919>
- <sup>147</sup> Münch, B., Trtik, P., Marone, F., & Stampanoni, M. (2009). Stripe and ring artifact  
<sup>148</sup> removal with combined wavelet—fourier filtering. *Optics Express*, 17(10), 8567–8591.  
<sup>149</sup> <https://doi.org/10.1364/OE.17.008567>
- <sup>150</sup> Vo, N. T., Atwood, R. C., & Drakopoulos, M. (2018). Superior techniques for eliminating  
<sup>151</sup> ring artifacts in x-ray micro-tomography. *Optics Express*, 26(22), 28396–28412. <https://doi.org/10.1364/OE.26.028396>